

D599: Data Preparation and Exploration Task 1

Staphon Smith

Part I: Data Profiling

A1a: General Characteristics/Profile Data:

The Employee Turnover Dataset given contains 10,200 Rows and 16 Columns. The data types that are included are: Integers, Floats and Objects.

A1b: Variable Data Type and Subtype and A1c: Observable Values

I used the attribute `.dtype` to give me a list of the datatypes of each column and this was the result that was given:

```
EmployeeNumber      int64
Age                  int64
Tenure               int64
Turnover             object
HourlyRate           object
HoursWeekly          int64
CompensationType     object
AnnualSalary         float64
DrivingCommuterDistance int64
JobRoleArea          object
Gender               object
MaritalStatus        object
NumCompaniesPreviouslyWorked float64
AnnualProfessionalDevHrs float64
PaycheckMethod       object
TextMessageOptIn     object
dtype: object
```

As defined by course materials/datacamp videos and the given above `.dtype` attributes, here are the datatypes, subtypes and first three observable values per column:

1. Employee Number
 - a. Data type: Categorical - (This is a unique identifier and should be treated as a categorical variable.)
 - b. Subtype: Identifier (Nominal)
 - c. Observable Examples: 1, 2, 3
2. Age
 - a. Data type: Numeric
 - b. Subtype: Float64
 - c. Observable Examples: 28,33,22
3. Tenure
 - a. Data type: Numeric
 - b. Subtype: Int64
 - c. Observable Examples: 6,2,1
4. Turnover

- a. Data type: Text/String
 - b. Subtype: Object
 - c. Observable Examples: Yes, No
- 5. Hourly Rate
 - a. Data type: Numeric
 - b. Subtype: Float64
 - c. Observable Examples: 24.37,22.52, 88.77
- 6. Hours Weekly
 - a. Data type: Numeric
 - b. Subtype: Int64
 - c. Observable Examples:40
- 7. Compensation Type
 - a. Data type:Text/String
 - b. Subtype: Object
 - c. Observable Examples: Salary
- 8. Annual Salary
 - a. Data type: Numeric
 - b. Subtype: Int64
 - c. Observable Examples: ,50689.6,46841.6, 284641.6
- 9. Driving Commuter Distance
 - a. Data type: Numeric
 - b. Subtype: Int64
 - c. Observable Examples: 89,35, 12
- 10. Job Role Area
 - a. Data type:Text/String
 - b. Subtype: Object
 - c. Observable Examples: Research, Information Technology, Sales
- 11. Gender
 - a. Data type:Text/String
 - b. Subtype: Object
 - c. Observable Examples: Female, Prefer Not to Answer, Male
- 12. Marital Status
 - a. Data type:Text/String
 - b. Subtype: Object
 - c. Observable Examples:Married, Single, Divorced
- 13. Number of Companies Previously Worked
 - a. Data type: Numeric
 - b. Subtype: Float64
 - c. Observable Examples: 3, 6, 1
- 14. Annual Professional Dev Hours
 - a. Data type:Numeric
 - b. Subtype: Float64
 - c. Observable Examples: 7,8,N/A
- 15. Paycheck Method
 - a. Data type:Text/String
 - b. Subtype: Object
 - c. Observable Examples: Mail Check, Mailed Check, Direct Deposit
- 16. Text Message Opt In
 - a. Data type:Text/String

- b. Subtype: Object
- c. Observable Examples: Yes, N/A, No

Part II: Data Cleaning and Plan

B1: Dataset Quality Issues and B2: List of Quality Issues

Import data into Dataframe:

```
import pandas as pd

file_path = r"C:\Users\StaphonSmith\Desktop\Employee Turnover Dataset (1).csv"

df=pd.read_csv(file_path)
```

Issue #1 - Duplicate Entries

After importing the data into the data frame I started with the first issue, duplicate entries. I used this code below and found there were 99 Duplicate Rows.

```
#Look for duplicated data
num_duplicates = df.duplicated().sum()
print(f'number of duplicate rows: {num_duplicates}')
```

```
number of duplicate rows: 99
```

Next, I printed the duplicate rows for quick inspection and then confirmed these 99 duplicates needed to be cleaned from the dataframe.

```
duplicate_rows = df[df.duplicated()]
print(duplicate_rows)
```

	EmployeeNumber	Age	Tenure	Turnover	HourlyRate	HoursWeekly	\
10100	1	28	6	Yes	\$24.37	40	
10101	2	33	2	Yes	\$24.37	40	
10102	3	22	1	No	\$22.52	40	
10103	4	23	1	No	\$22.52	40	
10104	5	40	6	No	\$88.77	40	
...	
10194	95	48	13	Yes	\$85.40	40	
10195	96	54	17	No	\$85.40	40	
10196	97	44	6	No	\$71.90	40	
10197	98	58	19	No	\$71.90	40	
10198	99	48	17	Yes	\$71.33	40	

	CompensationType	AnnualSalary	DrivingCommuterDistance	\
10100	Salary	50689.6	89	
10101	Salary	50689.6	89	
10102	Salary	46841.6	35	
10103	Salary	46841.6	35	
10104	Salary	284641.6	12	
...	
10194	Salary	177632.0	31	
10195	Salary	177632.0	31	
10196	Salary	149552.0	32	
10197	Salary	149552.0	32	
10198	Salary	148075.2	50	

	JobRoleArea	Gender	MaritalStatus	\
10100	Research	Female	Married	
10101	Research	Female	Married	

[99 rows x 16 columns]

Issue #2 - Missing Values

The next issue at hand was missing values. I used “.isnull” in my python code then summed up the totals to figure out which specific columns have missing values. Below is the code along with the data received.

```
#Look for missing data
print(df.isnull().sum())
```

EmployeeNumber	0
Age	0
Tenure	0
Turnover	0
HourlyRate	0
HoursWeekly	0
CompensationType	0
AnnualSalary	0
DrivingCommuterDistance	0
JobRoleArea	0
Gender	0
MaritalStatus	0
NumCompaniesPreviouslyWorked	665
AnnualProfessionalDevHrs	1969
PaycheckMethod	0
TextMessageOptIn	2266
dtype: int64	

3 of the 16 columns have more than one cell missing data, these will need to be evaluated and I will create a solution to resolve this.

Issues #3,4 and 5 - Inconsistent Entries, Formatting Errors, and Outliers

To streamline the process, I decided to check for formatting errors, inconsistencies, and outliers all at once. By examining the data distribution, it became easier to spot unusual entries, incorrect formatting, or values that didn't align with expected patterns. You can then find outliers by looking for extreme values. Obvious errors like negative distances were removed. For high but plausible values (e.g., long commute distances or large salaries), I evaluated them with domain context. I retained those that seemed realistic and removed only the most extreme, likely erroneous entries to avoid skewing the analysis. Below is a breakdown of the 16 columns I reviewed, along with the methods I used to detect and address these issues. I cleaned the data column by column, so in some cases, fixing earlier columns helped resolve problems in the later ones.

In the following Subsections I will elaborate in detail on the Code used in python and my findings (Screenshots included):

Employee Number:

Code:

```
#Check EmployeeNumber column for abnormalities  
EmployeeNumber_counts = df_cleaned['EmployeeNumber'].value_counts()  
print(EmployeeNumber_counts)
```

```
EmployeeNumber  
1            1  
6738         1  
6731         1  
6732         1  
6733         1  
..  
3367         1  
3368         1  
3369         1  
3370         1  
10100        1  
Name: count, Length: 10100, dtype: int64
```

Findings: Every Employee Number is matched properly and nothing to adjust.

Age:

Code:

```
#Check Age column for abnormalities  
Age_counts = df_cleaned['Age'].value_counts()  
print(Age_counts)
```

```
Age  
39    444  
37    422  
36    403  
38    382  
40    375  
43    317  
44    315  
48    311  
46    305  
56    303  
42    300  
41    299  
60    296  
58    296  
47    295  
54    293  
61    293  
59    290  
53    289  
45    281  
57    279  
49    278  
51    278  
52    266  
50    265  
55    243  
32    213  
30    202  
34    201
```

Findings: There were no abnormalities or outliers found in the age column.

Tenure:

Code:

```
Tenure_counts = df_cleaned['Tenure'].value_counts()
print(Tenure_counts)
```

```
Tenure
1      851
10     737
5      733
7      728
6      719
8      703
9      679
3      609
2      452
4      447
14     375
15     360
20     356
19     349
13     348
16     341
18     334
11     329
12     327
17     323
Name: count, dtype: int64
```

Findings: Tenure had no abnormalities.

Turnover:

Code:

```
#Check Turnover for abnormalities
Turnover_counts = df_cleaned['Turnover'].value_counts()
print(Turnover_counts)
```

```
Turnover
No      5456
Yes     4644
Name: count, dtype: int64
```

Findings: No abnormalities in Turnover.

Hourly Rate:

Code:

```
# I found there was an extra space after "Hourly Rate" so i used the code below so this will not happen again.
```

```
df_cleaned.columns = df_cleaned.columns.str.strip()
```

```
#Listing unique entries in HourlyRate
```

```
HourlyRate_counts = df_cleaned['HourlyRate'].value_counts()
```

```
# Display the counts
```

```
print("HourlyRate Counts:")
```

```
print(HourlyRate_counts)
```

```
HourlyRate  Counts:
```

```
HourlyRate
```

```
$34.28      11
```

```
$31.28      10
```

```
$33.66      10
```

```
$28.83       9
```

```
$33.06       9
```

```
..
```

```
$28.37       1
```

```
$56.02       1
```

```
$89.43       1
```

```
$88.05       1
```

```
$93.05       1
```

```
Name: count, Length: 5244, dtype: int64
```

Findings: there was a space after HourlyRate so i used the code stated above to remove that and avoid any issues further. Other than that there were no outliers or unusual data points.

Hours Weekly:

Code:

```
#Check HoursWeekly for abnormalities
```

```
HoursWeekly_counts = df_cleaned['HoursWeekly'].value_counts()
```

```
print(HoursWeekly_counts)
```

```
HoursWeekly
```

```
40      10100
```

```
Name: count, dtype: int64
```

Findings: Very standard 40 hour work weeks were found with no abnormalities.

Compensation Type:

Code:

```
#CompensationType for Errors/abnormalities
CompensationType_counts = df_cleaned['CompensationType'].value_counts()
print(CompensationType_counts)
```

```
CompensationType
Salary      10100
Name: count, dtype: int64
```

Findings: No issues found with Compensation Type.

Annual Salary:

Code:

```
#AnnualSalary for abnormalities
AnnualSalary_counts = df_cleaned['AnnualSalary'].value_counts()
print(AnnualSalary_counts)
```

```
AnnualSalary
76294.4      9
64896.0      8
54350.4      7
53414.4      7
67392.0      7
..
49254.4      1
156707.2     1
37086.4      1
202571.2     1
333544.0     1
Name: count, Length: 5538, dtype: int64
```

Findings: No issues with Annual Salary, given a normal company has wide ranges in pay for certain employees.

Driving Commuting Distance:

Code:

```
#Check DrivingCommuterDistance for abnormalities
```

```
DrivingCommuterDistance_counts = df_cleaned['DrivingCommuterDistance'].value_counts()  
print(DrivingCommuterDistance_counts)
```


```
DrivingCommuterDistance
```

```
33      184  
250     177  
28      157  
27      125  
42      113
```

```
...  
-125     19  
-275     17  
125      17  
910       4  
950       2
```

```
Name: count, Length: 120, dtype: int64
```

```
#there are a lot of outliers and commutes that are unfeasible. i would have taken out 250+ miles however there is many pec
```



```
# Remove rows where DrivingCommuterDistance is greater than 250 and less than 0
```

```
df_cleaned = df_cleaned[(df_cleaned['DrivingCommuterDistance'] <= 250) & (df_cleaned['DrivingCommuterDistance'] >= 0)]
```

```
# Remove rows where DrivingCommuterDistance is greater than 250 and less than 0
```

```
df_cleaned = df_cleaned[(df_cleaned['DrivingCommuterDistance'] <= 250) & (df_cleaned['DrivingCommuterDistance'] >= 0)]
```

```
#Reused prior code to see if outliers have been deleted.
```

```
DrivingCommuterDistance_counts = df_cleaned['DrivingCommuterDistance'].value_counts()  
print(DrivingCommuterDistance_counts)
```

```
DrivingCommuterDistance
```

```
33      184  
250     177  
28      157  
27      125  
42      113
```

```
...  
90      62  
98      58  
96      57  
91      51  
125     17
```

```
Name: count, Length: 100, dtype: int64
```

Findings: Found there were a few outliers that had negative distances. Obvious errors like negative distances were removed. For high but plausible values (e.g., long commute distances or large salaries), I evaluated them with domain context. I retained those that seemed realistic and removed only the most extreme, likely erroneous entries to avoid skewing the analysis.

Job Role Area:

Code:

```
#Listing unique entries in JobRoleArea
JobRoleArea_counts = df_cleaned['JobRoleArea'].value_counts()
```

```
# Display the counts
print("JobRoleArea:")
print(JobRoleArea_counts)
```

```
JobRoleArea:
JobRoleArea
Research          1733
Sales             1718
Marketing          943
Manufacturing     895
Healthcare        890
Laboratory        870
Human Resources   765
Information Technology 736
InformationTechnology 73
HumanResources    43
Information_Technology 36
Human_Resources   27
Name: count, dtype: int64
```

Findings: Nothing abnormal.

Gender:

Code:

```
#Listing unique entries in Gender
Gender_counts = df_cleaned['Gender'].value_counts()
```

```
# Display the counts
print(Gender_counts)
```

```
Gender
Female          4964
Male            3640
Prefer Not to Answer 125
Name: count, dtype: int64
```

Findings: Nothing Abnormal.

Marital Status:

Code:

```
#Listing unique entries in MaritalStatus
MaritalStatus_counts = df_cleaned['MaritalStatus'].value_counts()

# Display the counts
print("MaritalStatus:")
print(MaritalStatus_counts)
```

```
MaritalStatus:
MaritalStatus
Single      2939
Married     2924
Divorced    2866
Name: count, dtype: int64
```

Findings: Nothing abnormal and no outliers.

Number of Companies Previously Worked:

Code:

```
#Listing unique entries in NumCompaniesWorked
NumCompaniesPreviouslyWorked_counts = df_cleaned['NumCompaniesPreviouslyWorked'].value_counts()
print("NumCompaniesPreviouslyWorked:")
print(NumCompaniesPreviouslyWorked_counts)
```

```
NumCompaniesPreviouslyWorked:
NumCompaniesPreviouslyWorked
1.0      1289
2.0      1263
3.0      1259
6.0       914
4.0       873
5.0       846
7.0       588
8.0       565
9.0       561
Name: count, dtype: int64
```

Findings: Nothing Abnormal.

Annual Prof. Dev Hours:

Code:

```
#Listing unique entries in AnnualProfessionalDevHrs
AnnualProfessionalDevHrs_counts = df_cleaned['AnnualProfessionalDevHrs'].value_counts()
print("AnnualProfessionalDevHrs:")
print(AnnualProfessionalDevHrs_counts)
```

AnnualProfessionalDevHrs:

AnnualProfessionalDevHrs

14.0 368

10.0 368

23.0 365

25.0 363

13.0 352

5.0 346

6.0 344

22.0 340

12.0 339

9.0 334

19.0 332

7.0 330

18.0 329

15.0 329

17.0 326

11.0 325

8.0 324

24.0 317

21.0 314

20.0 310

16.0 305

Name: count, dtype: int64

Findings: Nothing Abnormal.

Pay Check Method:

Code:

```
PaycheckMethod_counts = df_cleaned['PaycheckMethod'].value_counts()
print("PaycheckMethod:")
print(PaycheckMethod_counts)
```

PaycheckMethod:

PaycheckMethod

Mail Check 4248

Mailed Check 2103

DirectDeposit 861

Direct_Deposit 799

Mail_Check 482

Direct Deposit 193

MailedCheck 43

Name: count, dtype: int64

#there is some errors in just putting data into two categories so I will fix that below.

```

# Clean and standardize the PaycheckMethod column
df_cleaned.loc[:, 'PaycheckMethod'] = df_cleaned['PaycheckMethod'].replace({
    'direct deposit': 'Direct Deposit',
    'check': 'Check'
})

# Map cleaned entries to standard categories
df_cleaned.loc[:, 'PaycheckMethod'] = df_cleaned['PaycheckMethod'].replace({
    'direct deposit': 'Direct Deposit',
    'check': 'Check',
    'Check': 'Check',
    'Direct deposit': 'Direct Deposit'
})

# Check result
print(df_cleaned['PaycheckMethod'].value_counts())

```

```

PaycheckMethod
Mailed Check      6876
Direct Deposit    1853
Name: count, dtype: int64

```

Findings: There were some issues with the spacing however no real issues. To clean up the data a bit and get rid of inconsistencies I grouped them properly into “Mailed Check” and “Direct Deposit” as shown above.

Text Messages Opt In:

Code:

```

TextMessageOptIn_counts = df_cleaned['TextMessageOptIn'].value_counts()
print("TextMessageOptIn:")
print(TextMessageOptIn_counts)

```

```

TextMessageOptIn:
TextMessageOptIn
Yes      6285
No       455
Name: count, dtype: int64

```

Findings: Nothing Abnormal.

C1:DATASET MODIFICATION and C2:DATA CLEANING TECHNIQUES

I will outline the issues identified in the dataset, the code used to address them, and the verification steps taken to confirm the corrections. Additionally, I will explain the importance and necessity of each cleaning method applied

Quality Issue #1 - Duplicate Entries

Problem: There were 99 duplicate rows to delete. Because they are duplicates, they are not useful.

Solution:

```
# Drop all duplicate rows
df_cleaned = df.drop_duplicates()
```

```
#Checking again for duplicated data to ensure they were deleted
num_duplicates = df_cleaned.duplicated().sum()
print(f'Number of duplicate rows: {num_duplicates}')
```

```
Number of duplicate rows: 0
```

Quality Issue #2 - Missing Values

Problem: 3 of the 16 columns have at least one cell of missing data.

The solution: Since the 3 columns that are missing values are plausible to have a 0 as an accurate measuring and to keep the bulk of that data i chose to fill the N/A's with 0 to maintain the data. As for Text Message Opt In i chose to change the N/A's to no as that would be a default answer.


```
# Fill missing values with 0 for numeric columns
df['NumCompaniesPreviouslyWorked'] = df['NumCompaniesPreviouslyWorked'].fillna(0)
df['AnnualProfessionalDevHrs'] = df['AnnualProfessionalDevHrs'].fillna(0)

# Fill missing values with 'No' for categorical column
df['TextMessageOptIn'] = df['TextMessageOptIn'].fillna('No')

# Confirm no missing values remain
print(df.isnull().sum())
```

```
EmployeeNumber      0
Age                 0
Tenure              0
Turnover            0
HourlyRate          0
HoursWeekly         0
CompensationType    0
AnnualSalary        0
DrivingCommuterDistance 0
JobRoleArea         0
Gender              0
MaritalStatus       0
NumCompaniesPreviouslyWorked 0
AnnualProfessionalDevHrs 0
PaycheckMethod      0
TextMessageOptIn    0
dtype: int64
```

```
# Check the number of rows before and after dropping missing values
print(f"Original DataFrame: {df.shape[0]} rows")
print(f"Cleaned DataFrame: {df_cleaned.shape[0]} rows")
```

```
Original DataFrame: 10199 rows
Cleaned DataFrame: 10100 rows
```

We kept 10,100 rows out of 10,199 after changing the missing values.

Quality Issue #3, 4, and 5 - Inconsistent Entries, Formatting Errors, Outliers

Here are the variables I found problems with after fixing the duplicate and missing values.

Hourly Rate:

Problem: There was an extra space that was causing formatting errors.

Solution:

```
# I found there was an extra space after "Hourly Rate" so i used the code below so this will not happen again
```

```
df_cleaned.columns = df_cleaned.columns.str.strip()
```

```
#Listing unique entries in HourlyRate
```

```
HourlyRate_counts = df_cleaned['HourlyRate'].value_counts()
```

```
# Display the counts
```

```
print("HourlyRate Counts:")
```

```
print(HourlyRate_counts)
```

```
HourlyRate  Counts:
```

```
HourlyRate
```

```
$34.28      11
```

```
$31.28      10
```

```
$33.66      10
```

```
$28.83       9
```

```
$33.06       9
```

```
..
```

```
$28.37       1
```

```
$56.02       1
```

```
$89.43       1
```

```
$88.05       1
```

```
$93.05       1
```

```
Name: count, Length: 5244, dtype: int64
```

Problematic data has been resolved.

Driving Commuter Distance:

Problem: there were outliers in this section where people were commuting over 250+ miles for work. also some negative mileage that would be impossible were also a problem.

Solution:

```
#Check DrivingCommuterDistance for abnormalities
```

```
DrivingCommuterDistance_counts = df_cleaned['DrivingCommuterDistance'].value_counts()
```

```
print(DrivingCommuterDistance_counts)
```

```
DrivingCommuterDistance
```

```
33      184
```

```
250     177
```

```
28      157
```

```
27      125
```

```
42      113
```

```
...
```

```
-125     19
```

```
-275     17
```

```
125      17
```

```
910       4
```

```
950       2
```

```
Name: count, Length: 120, dtype: int64
```

```
#there are a lot of outliers and commutes that are unfeasible. i would have taken out 250+ miles however there is many people who
```

```
# Remove rows where DrivingCommuterDistance is greater than 250 and less than 0
```

```
df_cleaned = df_cleaned[(df_cleaned['DrivingCommuterDistance'] <= 250) & (df_cleaned['DrivingCommuterDistance'] >= 0)]
```

I chose to clean those outliers from the data frame. Typically I would think 250 miles is quite far as well but didn't really qualify as an outlier as there were 177 others that also do that same commute.

```
#Reused prior code to see if outliers have been deleted.
DrivingCommuterDistance_counts = df_cleaned['DrivingCommuterDistance'].value_counts()
print(DrivingCommuterDistance_counts)
```

```
DrivingCommuterDistance
33      184
250     177
28      157
27      125
42      113
...
90       62
98       58
96       57
91       51
125      17
Name: count, Length: 100, dtype: int64
```

Problematic data has been resolved.

Paycheck Method

Problem: There were Inconsistent entries/formatting errors for the paycheck method.

Solution:

```
PaycheckMethod_counts = df_cleaned['PaycheckMethod'].value_counts()
print("PaycheckMethod:")
print(PaycheckMethod_counts)
```

```
PaycheckMethod:
PaycheckMethod
Mail Check      4248
Mailed Check    2103
DirectDeposit    861
Direct_Deposit   799
Mail_Check       482
Direct Deposit   193
MailedCheck       43
Name: count, dtype: int64
```

I chose to group them into 2 sections: "Direct Deposit" and "Mailed Check" for simplicity and better congruence.

```

# Clean and standardize the PaycheckMethod column
df_cleaned.loc[:, 'PaycheckMethod'] = df_cleaned['PaycheckMethod'].replace({
    'direct deposit': 'Direct Deposit',
    'check': 'Check'
})

# Map cleaned entries to standard categories
df_cleaned.loc[:, 'PaycheckMethod'] = df_cleaned['PaycheckMethod'].replace({
    'direct deposit': 'Direct Deposit',
    'check': 'Check',
    'Check': 'Check',
    'Direct deposit': 'Direct Deposit'
})

# Check result
print(df_cleaned['PaycheckMethod'].value_counts())

```

```

PaycheckMethod
Mailed Check      6876
Direct Deposit    1853
Name: count, dtype: int64

```

Problematic data has been Resolved.

C3: TECHNIQUE ADVANTAGES

I decided to remove the incomplete data instead of filling it in with averages or placeholders like null values. This approach ensures that the dataset is clean and consistent, which makes it easier to work with during analysis. We won't run into issues caused by missing or estimated values, and we can trust that all the information we're using is accurate and not based on assumptions.

C4: TECHNIQUE LIMITATIONS

By choosing to delete the missing data rather than substitute it with average values or nulls, we ended up with a smaller dataset. This might affect the overall outcome of our analysis, as those removed data points could have influenced the results. Additionally, if the missing values followed a specific pattern or were concentrated in a certain group, eliminating them entirely might introduce bias and leave out important trends.

D1-4: PANOPTO VIDEO & DOCUMENTS

These will be submitted together.

E: Sources

No sources were used besides WGU official course materials